

# Cluster Job Management with PBS

---

Andria

[andria.arisal@lipi.go.id](mailto:andria.arisal@lipi.go.id)

# Why?

---

- Resources are scarce
  - Users want to have them all
  - Administrator wants optimal resource utilization
- 
- So... jobs should be managed

# Outlines

---

- PBS Introduction
- Installation
- For Administrator
- For User

# PBS Introduction

---

# Portable Batch System

---

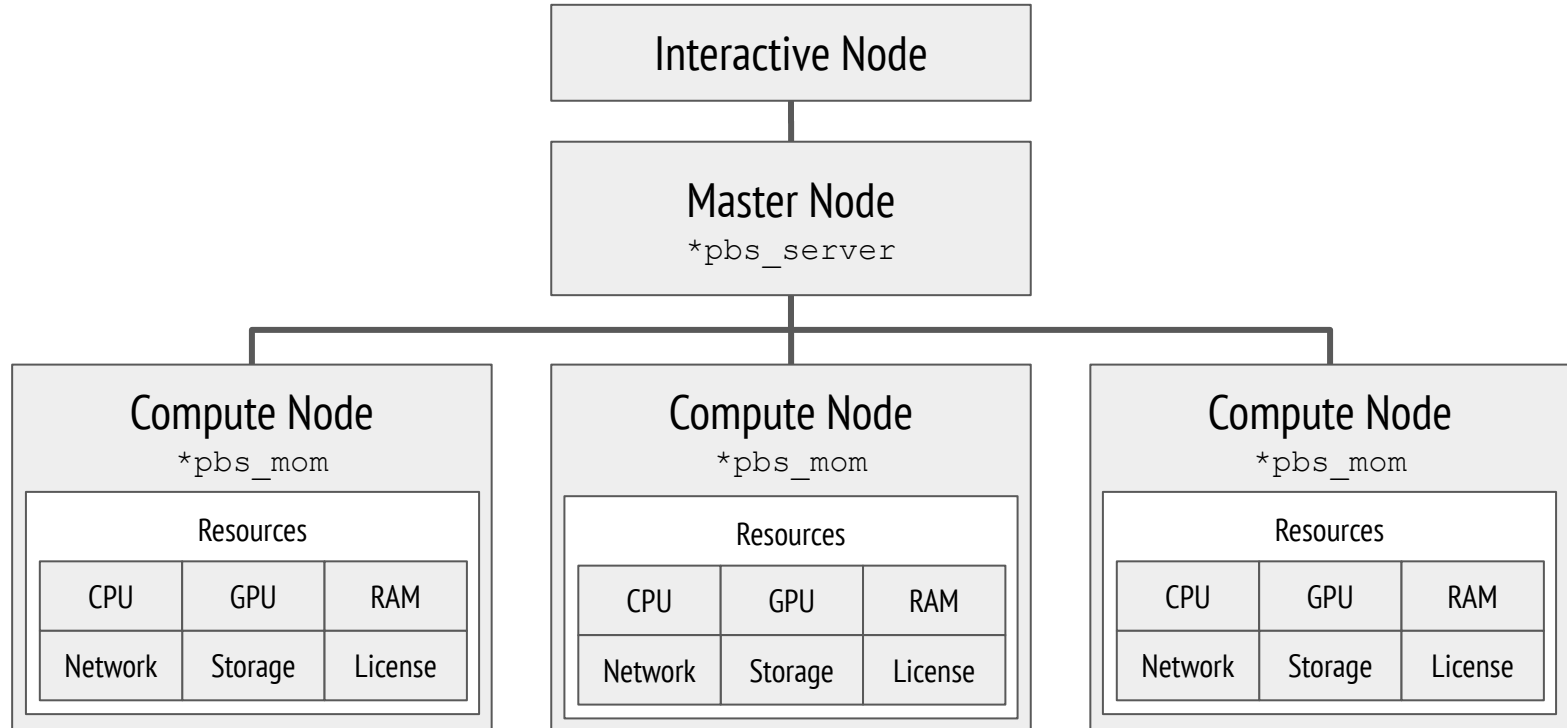
- Software for jobs scheduling
- Various fork
  - OpenPBS (no longer developed)
  - **TORQUE** (maintained by Adaptive Computing)
  - PBS Pro (commercially licensed by Altair Engineering)

# TORQUE

---

- Terascale Open-source Resource and QUEue Manager
- Resource manager for managing batch jobs in distributed compute nodes

# Concept



# Concept

Creation	Submission	Execution	Finalization
<pre>\$ vi job.pbs  #PBS -N myJob #PBS -S /bin/sh #PBS -l nodes=1:ppn=2,walltime=240:00:0 0 #PBS -M user@lapan.go.id #PBS -m ea source ~/.bashrc cd \$HOME/work/dir sh myBlast.sh -i -v</pre>	<pre>\$ qsub job.pbs</pre>	<pre>\$ qstat</pre>	<pre>\$ ls *.out \$ ls *.err</pre>



# MUNGE

---

- MUNGE (MUNGE Uid 'N' Gid Emporium)
- Authentication service for creating and validating credentials
- Security contract between users on the cluster
- Used by TORQUE

# Installation

---

# 0. Infrastructure

---

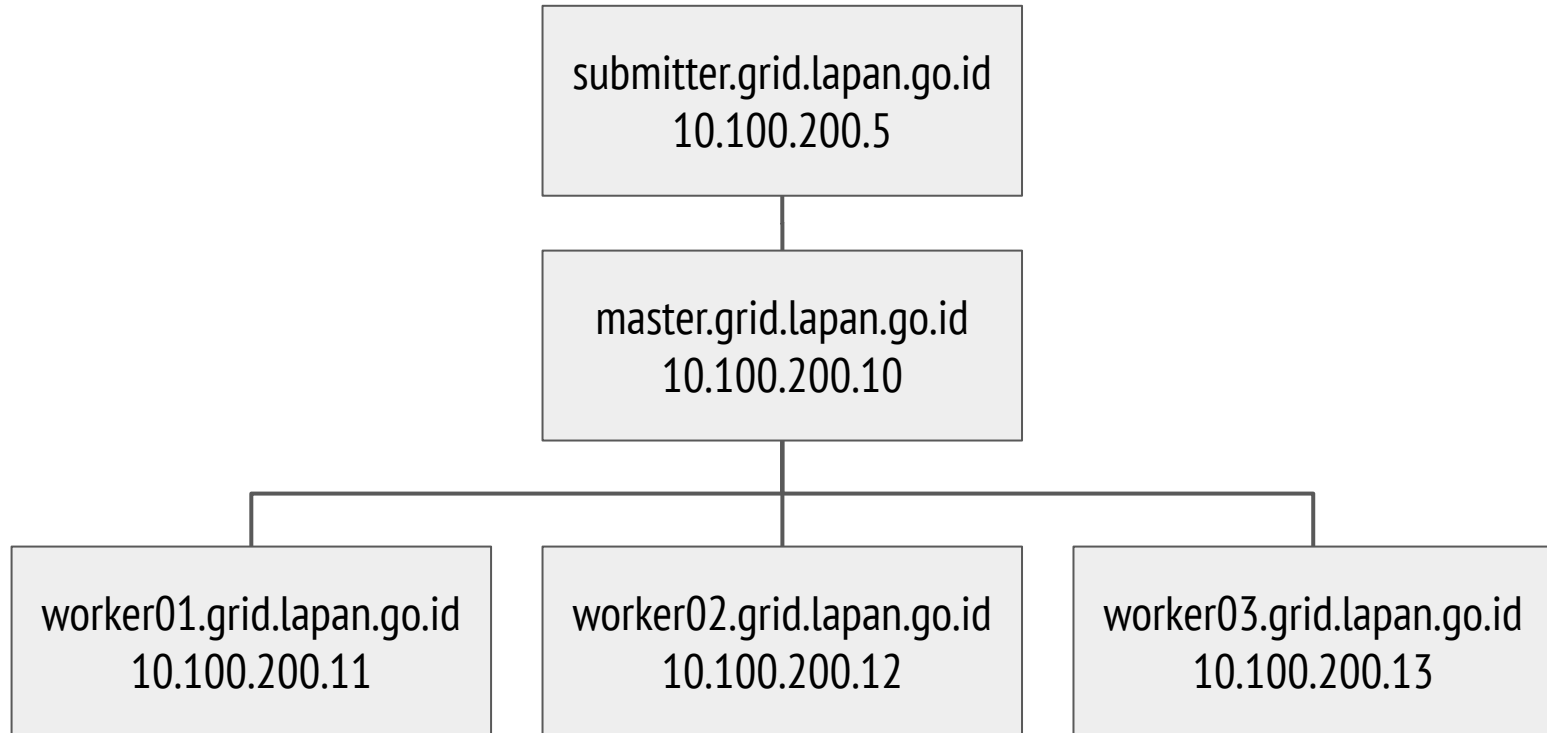
- Servers are installed with Operating System (CentOS 7)
- Network, firewall, passwordless-ssh, NFS, ... are configured

- **[root@\* ~]# cat /etc/hosts**

```
127.0.0.1      localhost.localdomain    localhost
::1           localhost.localdomain    localhost
10.100.200.5   submitter.grid.lapan.go.id submitter
10.100.200.10  master.grid.lapan.go.id  master
10.100.200.11  worker01.grid.lapan.go.id worker01
10.100.200.12  worker02.grid.lapan.go.id worker02
10.100.200.13  worker03.grid.lapan.go.id worker03
```

# 0. Infrastructure

---



# 1. PBS Master (installation)

---

- `[root@master ~]# yum install torque-server torque-scheduler`

...

Installed:

torque-scheduler.x86\_64 0:4.2.10-10.el7

torque-server.x86\_64 0:4.2.10-10.el7

Dependency Installed:

munge.x86\_64 0:0.5.11-3.el7

munge-libs.x86\_64 0:0.5.11-3.el7

torque.x86\_64 0:4.2.10-10.el7

torque-libs.x86\_64 0:4.2.10-10.el7

Complete!

# 1. PBS Master (MUNGE configuration and service)

---

- **[root@master ~]# /usr/sbin/create-munge-key**

Generating a pseudo-random key using /dev/urandom completed.

- **[root@master ~]# systemctl start munge**

- **[root@master ~]# systemctl enable munge**

- **[root@master ~]# systemctl status munge**

- munge.service - MUNGE authentication service

Loaded: loaded (/usr/lib/systemd/system/munge.service; enabled;  
vendor preset: disabled)

Active: active (running) since Mon 2019-01-28 08:55:33 WIB; 3s ago

Docs: man:munged(8)

Process: 5959 ExecStart=/usr/sbin/munged (code=exited,  
status=0/SUCCESS)

Main PID: 5961 (munged)

CGroup: /system.slice/munge.service

└─5961 /usr/sbin/munged

# 1. PBS Master (TORQUE configuration)

---

- `[root@master ~]# cat /var/lib/torque/server_name`  
`master`
- `[root@master ~]# cat /var/lib/torque/server_priv/nodes`  
`worker01.grid.lapan.go.id np=28 general`

# 1. PBS Master (TORQUE authentication service)

---

- `[root@master ~]# systemctl start trqauthd`
- `[root@master ~]# systemctl enable trqauthd`
- `[root@master ~]# systemctl status trqauthd`
  - `trqauthd.service - trqauthd`
    - `Loaded: loaded (/usr/lib/systemd/system/trqauthd.service; enabled; vendor preset: disabled)`
    - `Active: active (running) since Mon 2019-01-28 08:39:49 WIB; 17min ago`
    - `Process: 5398 ExecStart=/usr/sbin/trqauthd (code=exited, status=0/SUCCESS)`
    - `Main PID: 5399 (trqauthd)`
    - `CGroup: /system.slice/trqauthd.service`
      - `└─5399 /usr/sbin/trqauthd`



# 1. PBS Master (PBS service)

---

- `[root@master ~]# systemctl start pbs_server`
- `[root@master ~]# systemctl enable pbs_server`
- `[root@master ~]# systemctl status pbs_server`
  - `pbs_server.service - pbs-server`
    - Loaded: loaded (/usr/lib/systemd/system/pbs\_server.service; enabled; vendor preset: disabled)
    - Active: active (running) since Mon 2019-01-28 10:51:15 WIB; 1min 29s ago
    - Process: 2030 ExecStart=/usr/sbin/pbs\_server (code=exited, status=0/SUCCESS)
    - Main PID: 2032 (pbs\_server)
    - CGroup: /system.slice/pbs\_server.service
      - └─2032 /usr/sbin/pbs\_server

## 2. PBS Worker (installation)

---

- `[root@worker01 ~]# yum install torque-mom torque-client`

Installed:

```
torque-client.x86_64 0:4.2.10-10.el7  
torque-mom.x86_64 0:4.2.10-10.el7
```

Dependency Installed:

```
hwloc-libs.x86_64 0:1.11.8-4.el7  
munge.x86_64 0:0.5.11-3.el7  
munge-libs.x86_64 0:0.5.11-3.el7  
torque.x86_64 0:4.2.10-10.el7  
torque-libs.x86_64 0:4.2.10-10.el7
```

Complete!

## 2. PBS Worker (configuration)

---

- `[root@worker01 ~]# cat /var/lib/torque/server_name`  
`master.grid.lapan.go.id`
- `[root@worker01 ~]# cat /var/lib/torque/mom_priv/config`  
`# Configuration for pbs_mom.`  
`$pbsserver master.grid.lapan.go.id`  
`$logevent 0x1ff`
- `[root@master ~]# scp /etc/munge/munge.key worker01:/etc/munge/munge.key`
- `[root@worker01 ~]# chmod 400 /etc/munge/munge.key`

## 2. PBS Worker (MUNGE configuration and service)

---

- `[root@worker01 ~]# systemctl start munge`
- `[root@worker01 ~]# systemctl enable munge`
- `[root@worker01 ~]# systemctl status munge`
  - `munge.service - MUNGE authentication service`
    - `Loaded: loaded (/usr/lib/systemd/system/munge.service; enabled; vendor preset: disabled)`
    - `Active: active (running) since Mon 2019-01-28 08:55:33 WIB; 3s ago`
    - `Docs: man:munged(8)`
    - `Process: 5959 ExecStart=/usr/sbin/munged (code=exited, status=0/SUCCESS)`
    - `Main PID: 5961 (munged)`
    - `CGroup: /system.slice/munge.service`
      - `└─5961 /usr/sbin/munged`

## 2. PBS Worker (TORQUE authentication service)

---

- `[root@worker01 ~]# systemctl start trqauthd`
- `[root@worker01 ~]# systemctl enable trqauthd`
- `[root@worker01 ~]# systemctl status trqauthd`
  - `trqauthd.service - trqauthd`
    - `Loaded: loaded (/usr/lib/systemd/system/trqauthd.service; enabled; vendor preset: disabled)`
    - `Active: active (running) since Mon 2019-01-28 08:39:49 WIB; 17min ago`
    - `Process: 5398 ExecStart=/usr/sbin/trqauthd (code=exited, status=0/SUCCESS)`
    - `Main PID: 5399 (trqauthd)`
    - `CGroup: /system.slice/trqauthd.service`
      - `└─5399 /usr/sbin/trqauthd`

## 2. PBS Worker (PBS service)

---

- `[root@master ~]# systemctl start pbs_mom`
- `[root@master ~]# systemctl enable pbs_mom`
- `[root@master ~]# systemctl status pbs_mom`
  - `pbs_mom.service - pbs-mom`
    - `Loaded: loaded (/usr/lib/systemd/system/pbs_mom.service; enabled; vendor preset: disabled)`
    - `Active: active (running) since Mon 2019-01-28 10:35:21 WIB; 14min ago`
    - `Process: 2240 ExecStart=/usr/sbin/pbs_mom (code=exited, status=0/SUCCESS)`
    - `Main PID: 2241 (pbs_mom)`
    - `CGroup: /system.slice/pbs_mom.service`
      - `└─2241 /usr/sbin/pbs_mom`

### 3. PBS Submitter (installation)

---

- `[root@submitter ~]# yum install torque-client`

Installed:

```
torque-client.x86_64 0:4.2.10-10.el7
```

Dependency Installed:

```
hwloc-libs.x86_64 0:1.11.8-4.el7
```

```
munge.x86_64 0:0.5.11-3.el7
```

```
munge-libs.x86_64 0:0.5.11-3.el7
```

```
torque.x86_64 0:4.2.10-10.el7
```

```
torque-libs.x86_64 0:4.2.10-10.el7
```

Complete!

### 3. PBS Submitter (configuration)

---

- `[root@rumah ~]# cat /var/lib/torque/server_name`  
`master.grid.lapan.go.id`
- `[root@master ~]# scp /etc/munge/munge.key submitter:/etc/munge/munge.key`
- `[root@submitter ~]# chmod 400 /etc/munge/munge.key`



### 3. PBS Submitter (MUNGE configuration and service)

---

- `[root@submitter ~]# systemctl start munge`
- `[root@submitter ~]# systemctl enable munge`
- `[root@submitter ~]# systemctl status munge`
  - `munge.service - MUNGE authentication service`
    - `Loaded: loaded (/usr/lib/systemd/system/munge.service; enabled; vendor preset: disabled)`
    - `Active: active (running) since Mon 2019-01-28 08:55:33 WIB; 3s ago`
    - `Docs: man:munged(8)`
    - `Process: 5959 ExecStart=/usr/sbin/munged (code=exited, status=0/SUCCESS)`
    - `Main PID: 5961 (munged)`
    - `CGroup: /system.slice/munge.service`
      - `└─5961 /usr/sbin/munged`

### 3. PBS Submitter (TORQUE authentication service)

---

- `[root@submitter ~]# systemctl start trqauthd`
- `[root@submitter ~]# systemctl enable trqauthd`
- `[root@submitter ~]# systemctl status trqauthd`
  - `trqauthd.service - trqauthd`
    - `Loaded: loaded (/usr/lib/systemd/system/trqauthd.service; enabled; vendor preset: disabled)`
    - `Active: active (running) since Mon 2019-01-28 08:39:49 WIB; 17min ago`
    - `Process: 5398 ExecStart=/usr/sbin/trqauthd (code=exited, status=0/SUCCESS)`
    - `Main PID: 5399 (trqauthd)`
    - `CGroup: /system.slice/trqauthd.service`
      - `└─5399 /usr/sbin/trqauthd`

# for Administrator

---

# Creating Queue (queue definition)

---

- `[root@master ~]# vi qmgr.input`  
#  
# Create queues and set their attributes.  
#  
#  
# Create and define queue admin  
#  
create queue admin  
set queue admin queue\_type = Execution  
set queue admin enabled = True  
set queue admin started = True

# Creating Queue (queue definition)

---

- `[root@master ~]# vi qmgr.input`

...

```
create queue public
```

```
set queue public queue_type = Execution
```

```
set queue public max_user_queuable = 4
```

```
set queue public resources_max.nodes = 16
```

```
set queue public resources_max.cput = 48:00:00
```

```
set queue public resources_max.walltime = 72:00:00
```

```
set queue public resources_default.neednodes = general
```

```
set queue public resources_default.nodes = 4
```

```
set queue public resources_default.walltime = 24:00:00
```

```
set queue public max_user_run = 2
```

```
set queue public enabled = True
```

```
set queue public started = True
```

# Creating Queue (server attributes definition)

---

- `[root@master ~]# vi qmgr.input`  
...  
`set server scheduling = True`  
`set server acl_host_enable = False`  
`set server acl_hosts = master.grid.lapan.go.id`  
`set server acl_hosts += master`  
`set server managers = grid@mail.lapan.go.id`  
`set server managers += root@master.grid.lapan.go.id`  
`set server default_queue = default`  
`set server log_events = 511`  
`set server mail_from = adm`  
`set server scheduler_iteration = 600`  
`set server node_check_rate = 150`  
`set server tcp_timeout = 6`

# Creating Queue (server attributes definition)

---

- `[root@master ~]# vi qmgr.input`  
...  
`set server default_node = worker`  
`set server node_pack = False`  
`set server job_stat_rate = 45`  
`set server poll_jobs = True`  
`set server mom_job_sync = True`  
`set server mail_domain = never`  
`set server kill_delay = 10`  
`set server submit_hosts = submitter.grid.lapan.go.id`  
`set server allow_node_submit = True`  
`set server next_job_number = 7739`  
`set server moab_array_compatible = True`  
`set server nppcu = 1`

# Creating Queue (queue definition loading)

---

- `[root@master ~]# qmgr < qmgr.input`



# Monitoring Queue

---

- `[root@master ~]# qstat -f -Q`

...

Queue: public

queue\_type = Execution

max\_user\_queuable = 4

total\_jobs = 0

state\_count = Transit:0 Queued:0 Held:0 Waiting:0 Running:0 Exiting:0 Complete:0

resources\_max.cput = 48:00:00

resources\_max.nodes = 16

resources\_max.walltime = 72:00:00

resources\_default.neednodes = worker

resources\_default.nodes = 4

resources\_default.walltime = 24:00:00

mtime = 1548983368

resources\_assigned.nodect = 0

max\_user\_run = 2

enabled = True

started = True

# Controlling & Monitoring PBS

---

- `[root@master ~]# pbs_server [options]`
- `[root@master ~]# qterm`
- `[root@master ~]# qmgr`
  
- `[root@master ~]# pbsnode [options]`

# Controlling Queue

---

- `[root@master ~]# qstat`
  - `[root@master ~]# qhold`
  - `[root@master ~]# qrls`
  - `[root@master ~]# qalter`
  - `[root@master ~]# qdel`
- 
- `[root@master ~]# qstart`
  - `[root@master ~]# qstop`
  - `[root@master ~]# qrun`
  - `[root@master ~]# qselect`
  - `[root@master ~]# qorder`

# for User

---

# PBS Job script

---

- `[root@submitter ~]# vi myJob.pbs`

```
#!/bin/bash
```

```
#PBS -N myJob
```

```
#PBS -q public
```

```
#PBS -l nodes=1
```

```
#PBS -l walltime=1:00:00
```

```
#PBS -k oe
```

```
date
```

```
hostname
```

```
sleep 30
```

```
date
```

# Managing Job

---

- `[root@master ~]# qstat`
- `[root@master ~]# qhold`
- `[root@master ~]# qrls`
- `[root@master ~]# qalter`
- `[root@master ~]# qdel`